

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from sklearn.preprocessing import StandardScaler
```

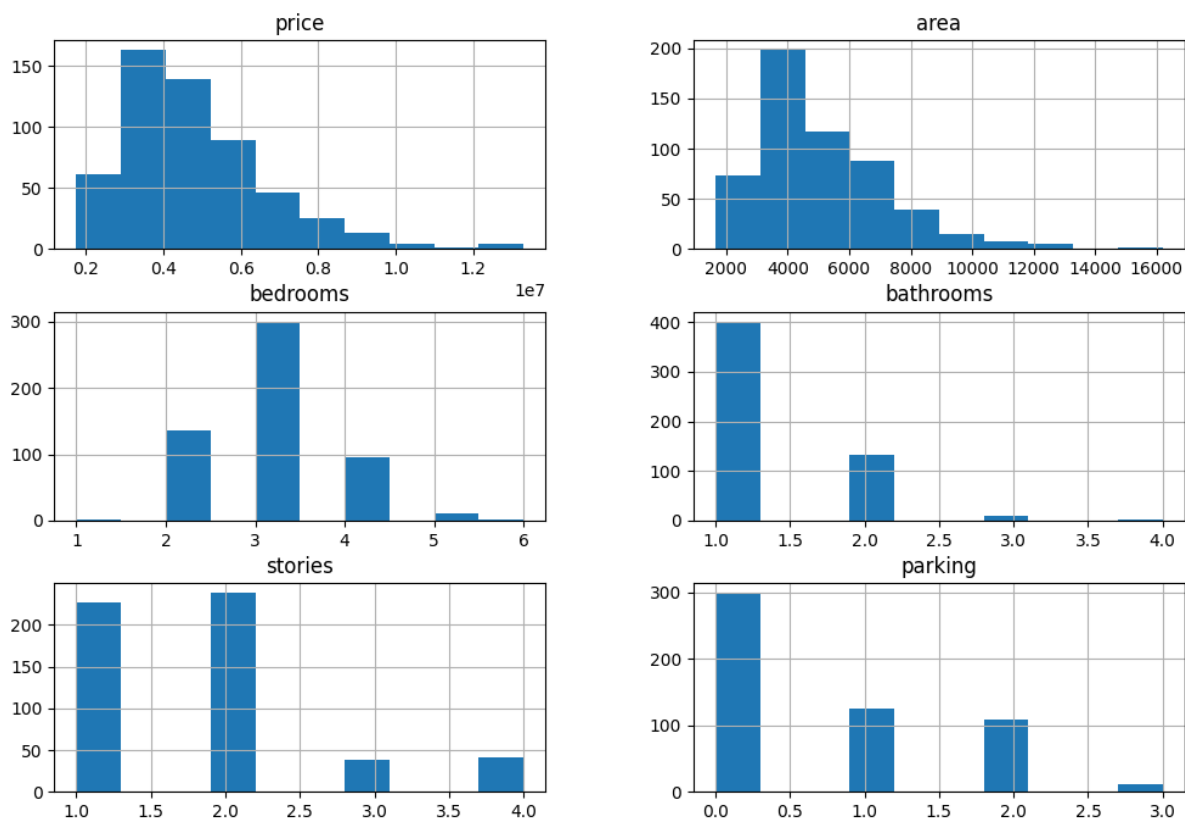
C:\Users\Tia\AppData\Local\Programs\Python\Python310\lib\site-packages\numpy_dist
 ribrator_init.py:30: UserWarning: loaded more than 1 DLL from .libs:
 C:\Users\Tia\AppData\Local\Programs\Python\Python310\lib\site-packages\numpy\.libs
 \libopenblas.FB5AE2TYXYH2IJRDKGQ3XBLKTF43H.gfortran-win_amd64.dll
 C:\Users\Tia\AppData\Local\Programs\Python\Python310\lib\site-packages\numpy\.libs
 \libopenblas64_v0.3.21-gcc_10_3_0.dll
 warnings.warn("loaded more than 1 DLL from .libs:")

In [2]:

```
# Step 1: Download and load the dataset
data = pd.read_csv('Housing.csv')
```

In [3]:

```
# Step 3: Perform visualizations
# Univariate Analysis
data.hist(figsize=(12, 8))
plt.show()
```

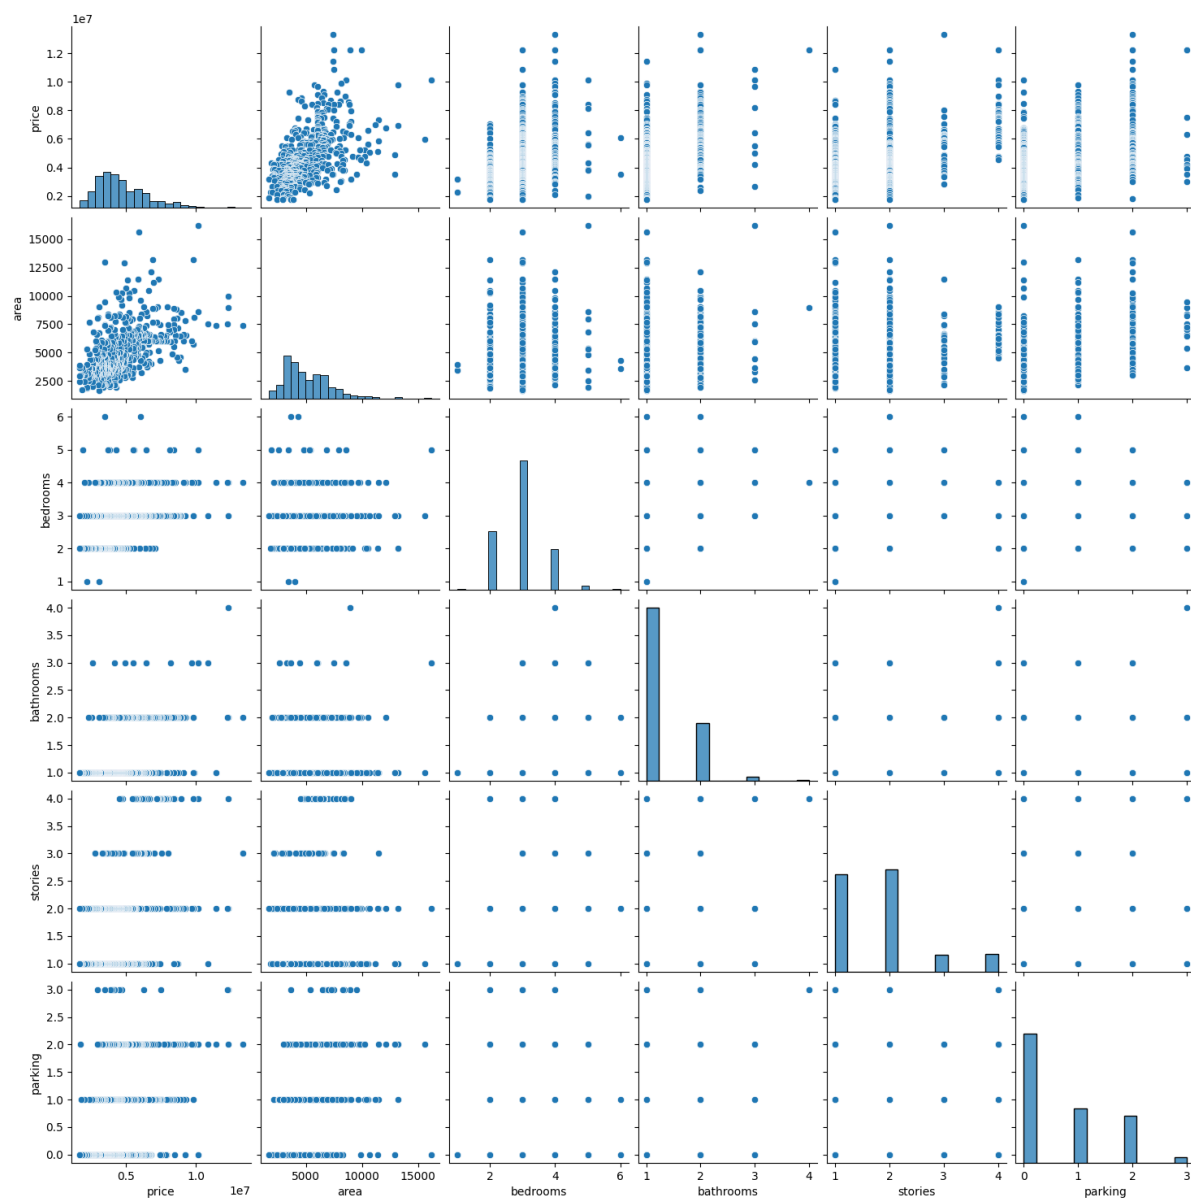


In [4]:

```
# Bi-Variate Analysis
```

```
sns.pairplot(data)
```

```
plt.show()
```



In [7]:

```
# Multi-Variate Analysis
```


```
numeric_cols = ['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'parking']
```

```
data_numeric = data[numeric_cols]
```

```
correlation_matrix = data_numeric.corr()
```

```
sns.heatmap(correlation_matrix, annot=True)
```

```
plt.show()
```



In [8]:

```
# Step 4: Perform descriptive statistics
statistics = data.describe()
print(statistics)
```

	price	area	bedrooms	bathrooms	stories
count	5.450000e+02	545.000000	545.000000	545.000000	545.000000
mean	4.766729e+06	5150.541284	2.965138	1.286239	1.805505
std	1.870440e+06	2170.141023	0.738064	0.502470	0.867492
min	1.750000e+06	1650.000000	1.000000	1.000000	1.000000
25%	3.430000e+06	3600.000000	2.000000	1.000000	1.000000
50%	4.340000e+06	4600.000000	3.000000	1.000000	2.000000
75%	5.740000e+06	6360.000000	3.000000	2.000000	2.000000
max	1.330000e+07	16200.000000	6.000000	4.000000	4.000000

	parking
count	545.000000
mean	0.693578
std	0.861586
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	3.000000

In [9]:

```
# Step 5: Check for missing values and deal with them
missing_values = data.isnull().sum()
print(missing_values)
```

```
price          0
area           0
bedrooms       0
bathrooms      0
stories        0
mainroad       0
guestroom      0
basement       0
hotwaterheating 0
airconditioning 0
parking        0
furnishingstatus 0
dtype: int64
```

In [11]:

```
# Step 7: Check for categorical columns and perform encoding
categorical_cols = ['mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'parking']
label_encoder = LabelEncoder()
for col in categorical_cols:
    data[col] = label_encoder.fit_transform(data[col])
```

In [13]:

```
# Perform one-hot encoding for categorical variables
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder

column_transformer = ColumnTransformer([('encoder', OneHotEncoder(), [11])], remainder='passthrough')
data = column_transformer.fit_transform(data)
```

In [14]:

```
# Step 8: Split the data into dependent and independent variables
X = data[:, 1:] # Independent variables
y = data[:, 0] # Dependent variable
```

In [15]:

```
# Step 9: Scale the independent variables
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

In [16]:

```
# Step 10: Split the data into training and testing
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

In [17]:

```
# Step 11: Build the Model (Linear Regression)
model = LinearRegression()
```

In [18]:

```
# Step 12: Train the Model
model.fit(X_train, y_train)
```

Out[18]:

```
▼ LinearRegression
LinearRegression()
```

In [19]:

```
# Step 13: Test the Model
y_pred = model.predict(X_test)
```

In [22]:

```
from sklearn.metrics import mean_squared_error, mean_absolute_error

# Step 14: Measure the performance using Metrics
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
mae = mean_absolute_error(y_test, y_pred)

print("Mean Squared Error:", mse)
print("Root Mean Squared Error:", rmse)
print("Mean Absolute Error:", mae)
```

```
Mean Squared Error: 7.602511275517291e-31
Root Mean Squared Error: 8.719238083409175e-16
Mean Absolute Error: 7.33358328192764e-16
```

In []: